

COMS31700 Design Verification: **Checking**

Kerstin Eder

(Acknowledgement: Avi Ziv from the IBM Research Labs in Haifa has kindly permitted the re-use of some of his slides.)

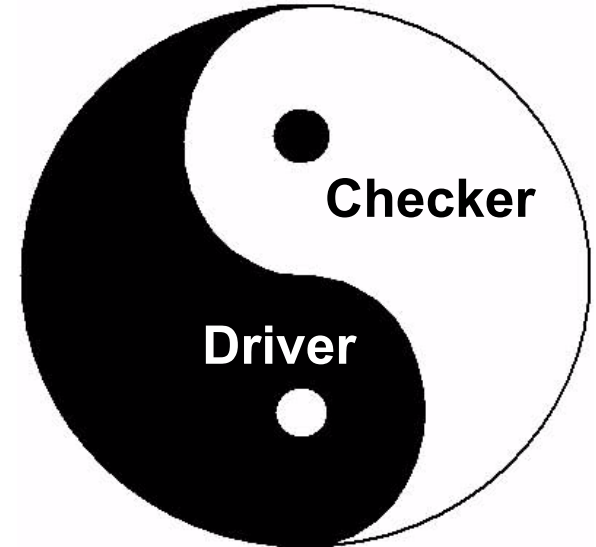
Checking - Outline

- Motivation
- Issues in checking
 - When to check
 - What to check
- Checking technologies
 - Reference models
 - Scoreboards
 - (Rule-based checking)
 - (Assertions)
 - Assertion-based verification (ABV) - later



The Yin-Yang of Verification

- Driving and checking are the yin and yang of verification
 - We cannot find bugs without creating the failing conditions.
 - We cannot find bugs without detecting the incorrect behavior.

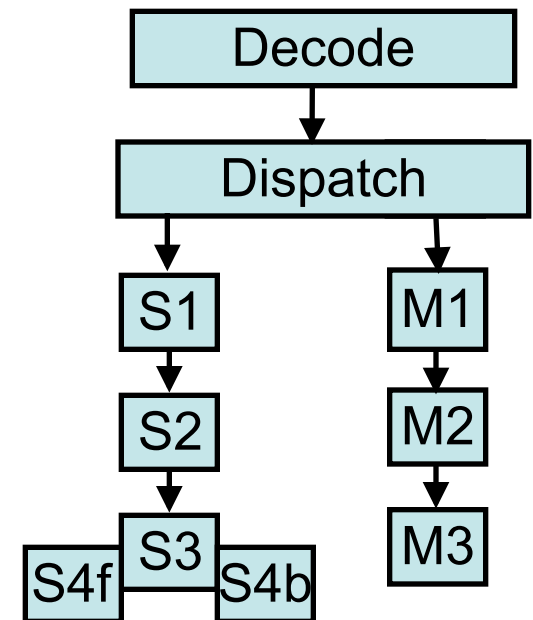


Ideal Checking

- In theory – detect deviation from expected behavior as soon as it happens and where it happens
 - No need to worry about “disappearing errors”
 - Easy to debug – the checker points to the bug
- This is not easy (even if we ignore many practical aspects) because in many cases we understand that something bad happened only in retrospect
 - Several “good” behaviors collide to create a bad behavior
- And what about the bugs we are not looking for?

“Good” Behavior Collision

- At cycle 1000 `fdiv F1, F2, F3` is dispatched to the M unit
 - It reaches stage M2 at cycle 1001
 - Its execution time is 60 cycles
- At cycle 1023 `fld F1, 100(G2)` is dispatched to the S unit
 - It reaches stage S2 at cycle 1024
- The data returns from the cache at cycle 1060
- At cycle 1061 the `fdiv` is ready to write
 - It moves to stage M3
- At cycle 1061 the `fld` is ready to write
 - It moves to stage S3
- Both instructions write to the same register together



“Good” Behavior Collision

- There are many possible causes for the problem

Checking: Practical Aspects

Consider:

- The cost of *implementation and maintenance* of checkers

vs

- the cost of *debugging* (without checkers).
- The cost of mistakes
 - Missed detection
 - We failed to detect a bug that was exposed by the stimuli.
 - False alarm
 - We mistakenly flagged a good behavior as bad.

Which is more expensive?

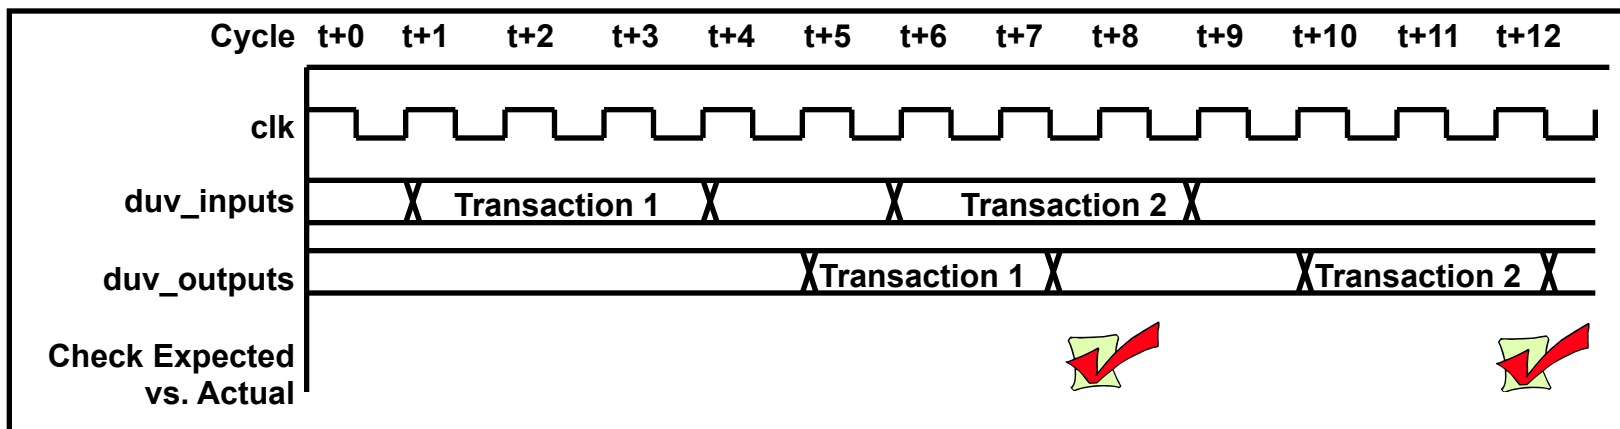
When to check

When to Check?

- Checking can be done at various stages of the verification job
 - During simulation
 - On-the-fly checking
 - At the end of simulation
 - End-of-test checking
 - After the verification job finishes
 - External checking
- Checking at each stage has its own advantages and disadvantages

On-the-fly Checking

- Checking is done **while the simulation is running**
- The DUV is continuously monitored to detect erroneous behavior



On-the-fly Checking

■ Advantages

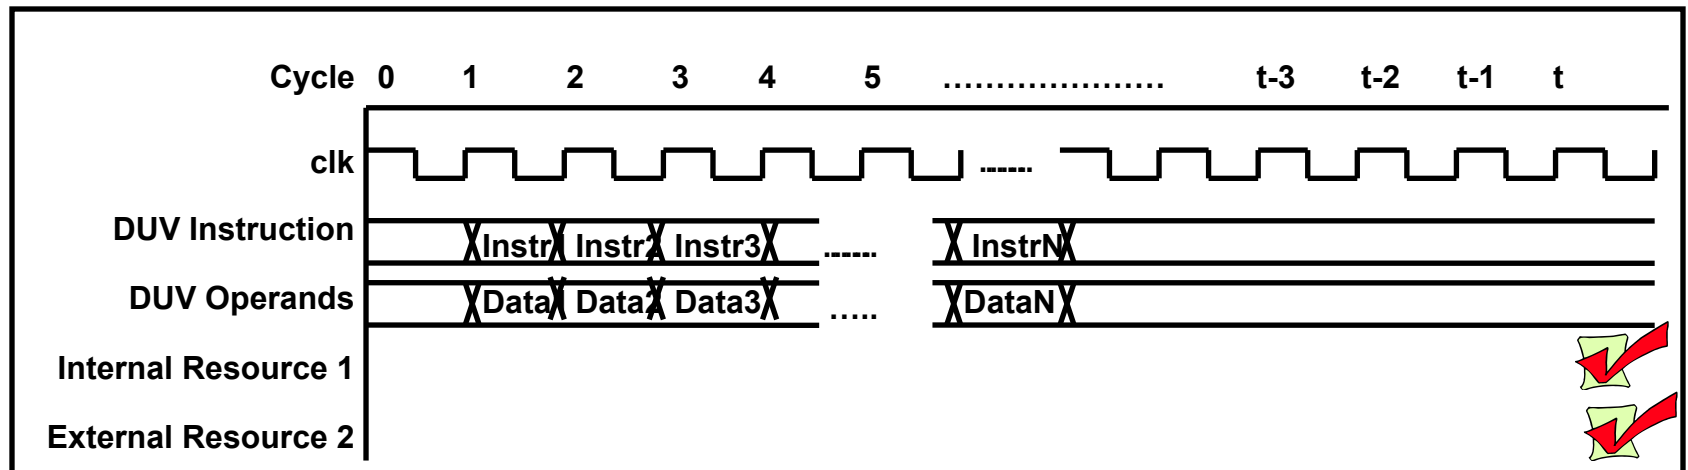
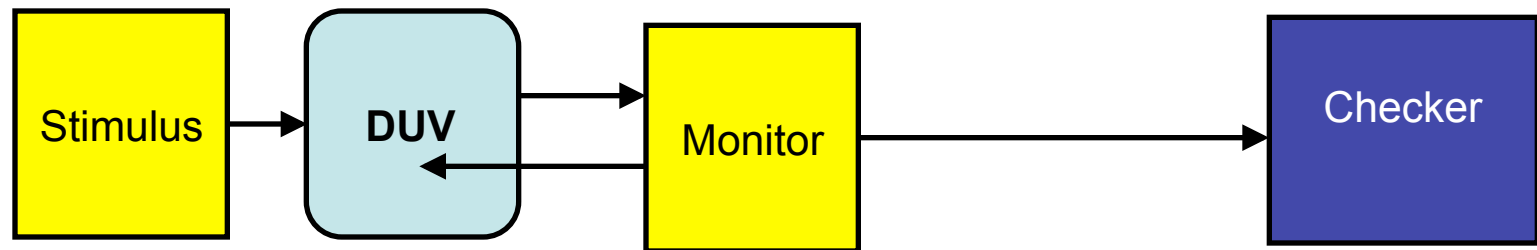
- Detection can be as close as possible (in time and space) to the bug source
- Can stop test as soon as bug occurs; no wasted simulation cycles
- Do not require large traces and external tools to do the checking

■ Disadvantages

- May **slow down simulation**
- Checking is limited to allowed time and space complexity
- Need to **plan the checking in advance**
 - To add a new checker, we need to rerun simulation

End-of-test Checking

- Checking is done **at the end of simulation**
- The checker checks the state of internal and external resources and makes sure that they are correct



End-of-test Checking

■ Disadvantages

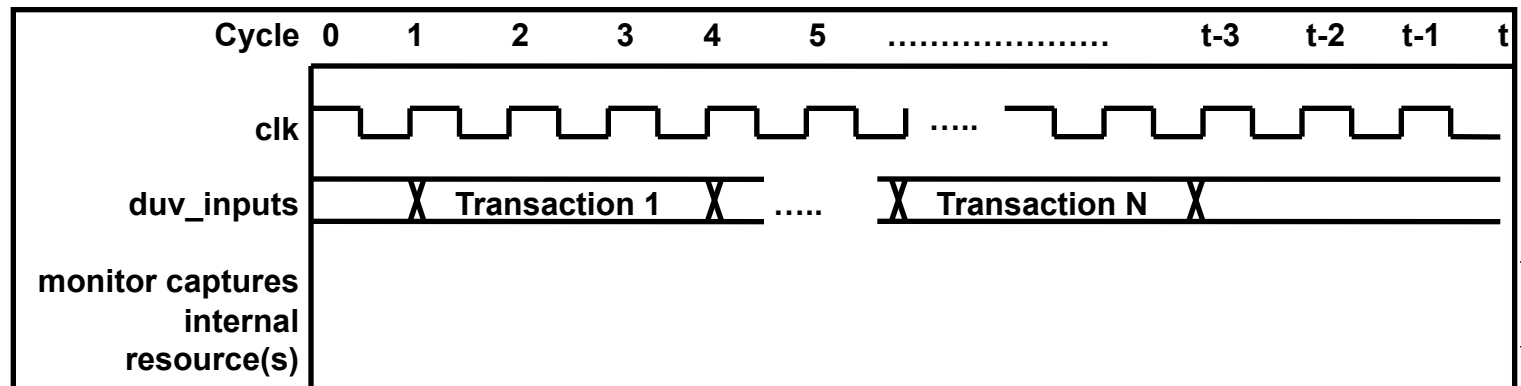
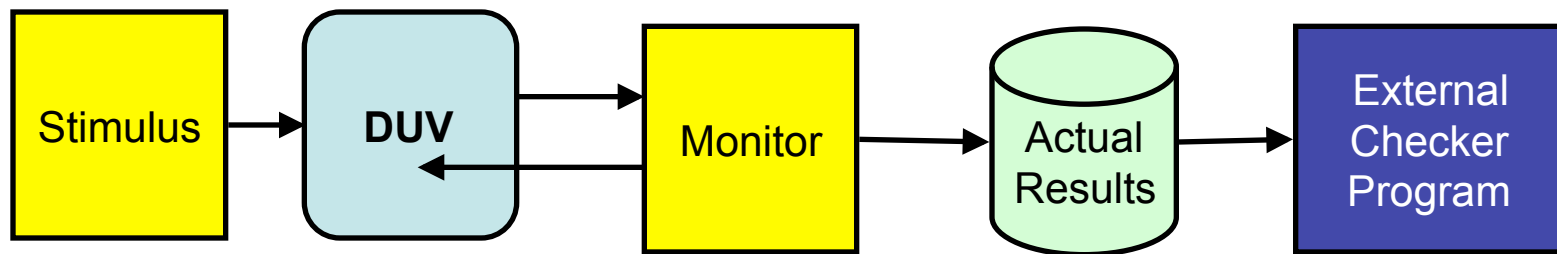
- Provides limited checking capabilities
 - Static look at the state of resources at the end of the test
- High probability of **masking bugs** by repeated writing to the resources
- **Hard to detect performance bugs**
 - Correct things are happening, but not at the right time
- Hard to correlate symptoms to bugs
 - **Hard to debug**

■ Advantages

- Simpler than other forms
 - May not require a deep understanding of the DUV
- Reduces probability of false alarms
 - (because bad effects may disappear)

External Checking (Monitors)

- **Monitors** keep internal resources' values and behaviors in trace files
- Checking is done by an external program that examines these files



External Checking

- External checking **separates the checking from the simulation**
 - We can perform any check we want without rerunning the simulation
 - As long as the data is in the trace files
 - We can perform more complicated checks
 - Use longer history, process events out-of-order
 - We can combine information coming from different sources
 - For example, different verification environments

In theory, external checking has all the powers of on-the-fly checking plus end-of-test checking - plus more

(Trace size and amount of traced facilities is a practical limitation.)

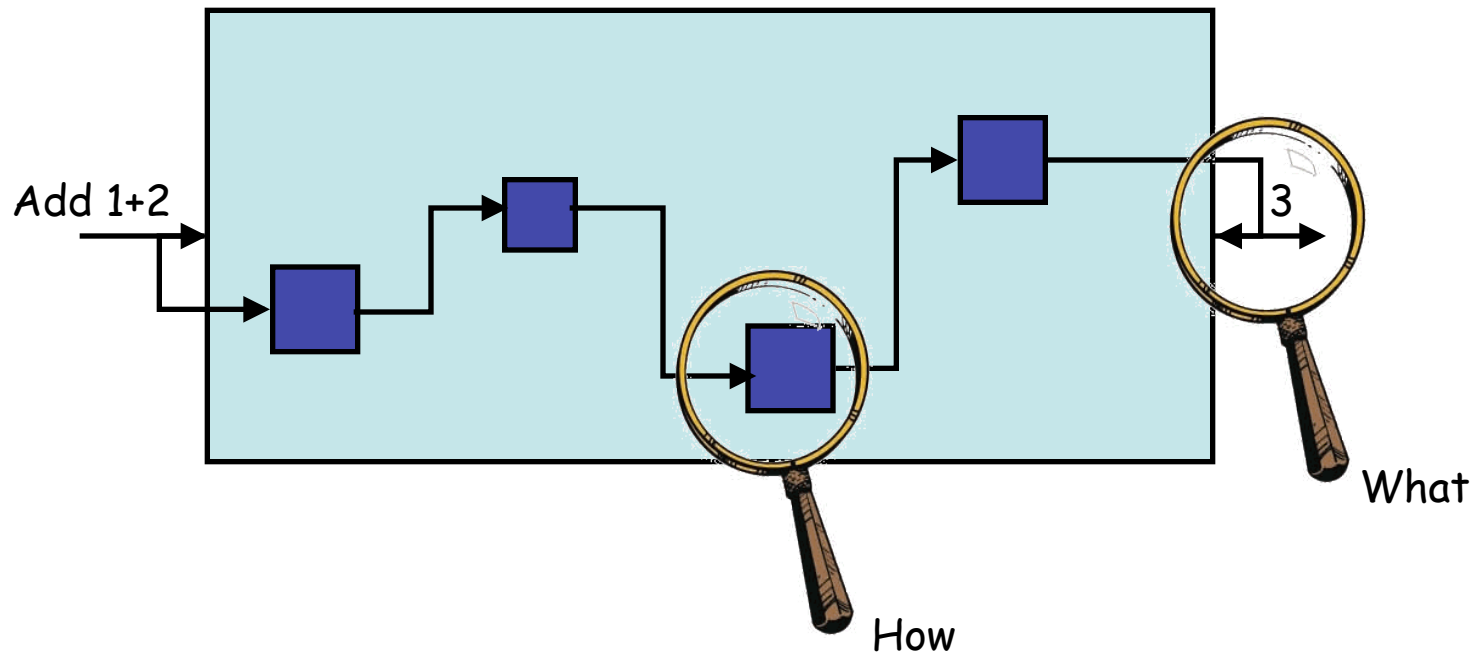
What to check

What to Check

- There are five main sources of checkers
 - The **inputs and outputs** of the design (specification)
 - The **architecture** of the design
 - The **microarchitecture** of the design
 - The **implementation** of the design
 - The **context** of the design
 - e.g. protocol compliance

(Slide from a previous lecture to remind us of where we can get inspiration for checkers from.)

Coarser Classification – The What and the How



Checking the What

- Check the **final outcome** of a behavior
 - Data oriented
 - But not limited to data
 - Usually based on higher level of abstraction
 - Checking is less tight
 - Requires less familiarity with the DUV
 - Low correlation between **failure**, the observed behavior that violates the spec, and **bugs/faults**, the root cause of the failure
 - Harder for debugging
 - Can find “unexpected” bugs

Checking the How

- Check *how* things are done internally
 - Control oriented
 - Usually at lower levels of abstraction
 - Closer to implementation
 - Tighter correlation between failure and bugs

Stimuli Generation and Checking

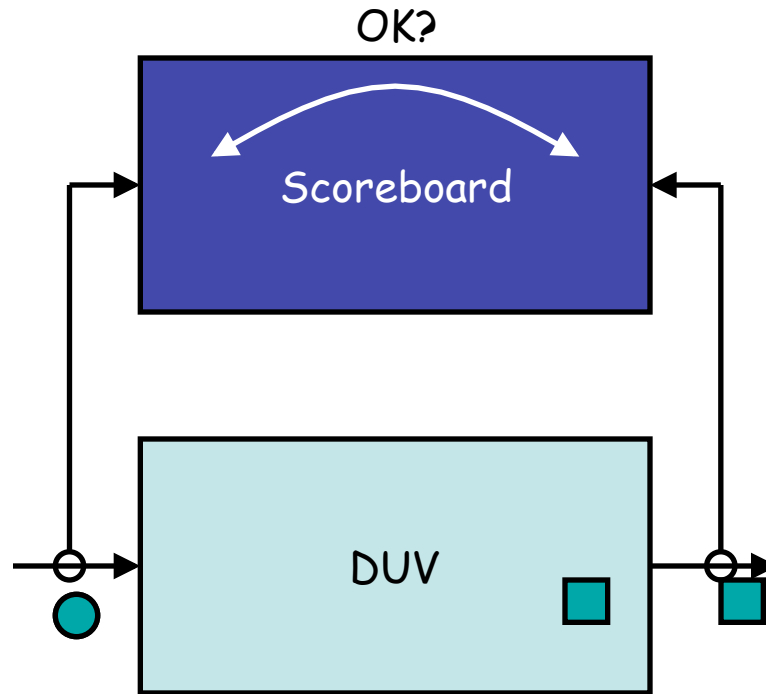
- In general, checking should be isolated from the stimuli generation
 - **Modularity** – ability to replace the stimuli generator
 - **Reusability** – ability to use the checkers at higher level of the design hierarchy
 - **Independence of Checking from Generation**
- Exceptions
 - Self-checking tests
 - Golden vectors
- *Can stimuli generation assist checking?*
 - The stimuli generation can assist checking by improving observability
 - Help transfer events from dark corners to the spotlight

Checking Technologies

Scoreboards

- Scoreboards are **smart data structures** that keep track of events in the DUV during simulation
- Usually, scoreboards are global
 - **One scoreboard per verification environment**
- Scoreboards are not checking mechanisms, but
 - The main purpose of using scoreboards is for checking
 - In practice, many checkers are implemented inside scoreboards
 - There are many typical checks that are done with scoreboards

Scoreboard Operation



Scoreboards Overview

- Scoreboards source information from
 - the inputs and outputs of the DUV, and
 - occasionally also from internal events in the DUV.
- **Scoreboards are very useful in dataflow designs**
 - Routers, cache designs, queues
- Types of checks enabled using a scoreboard:
 - Matching outputs with inputs
 - No loss of data
 - Detect inputs with no matching output.
 - No creation of data
 - Detect output with no matching input.
 - No unintended modification of data
 - Timing specification
 - Delay from input to output remains within specified limits.
 - Data order

Scoreboarding in e - 1

- Assume: DUV does not change order of packets.
 - Hence, first packet on scoreboard has to match received packet.

```
import packet_s;
unit scoreboard {
    !expected_packets : list of packet_s;
    add_packet(p_in : packet_s) is {
        expected_packets.add(p_in);
    };
    check_packet(p_out : packet_s) is {
        var diff : list of string;
        -- Compare physical fields of first packet on scb with p_out.
        -- Report up to 10 differences.
        diff = deep_compare_physical(expected_packets[0], p_out, 10);
        check that (diff.is_empty())
        else dut_error("`Packet not found on scoreboard'",
diff);
        -- If match was successful, continue.
        out("`Found received packet on scoreboard.'");
        expected_packets.delete(0);
    };
};
```

Scoreboarding in e - 2

Recording a packet on the scoreboard:

Extend driver such that

- When packet is driven into DUV call **add_packet** method of scoreboard.
 - Current packet is copied to scoreboard.
- It is useful to define an **event** that indicates when packet is being driven.

Checking for a packet on the scoreboard:

Extend receiver such that

- When a packet was received from DUV call **check_packet**.
 - Try to find the matching packet on scoreboard.
- It is useful to define an **event** that indicates when a packet is being received.

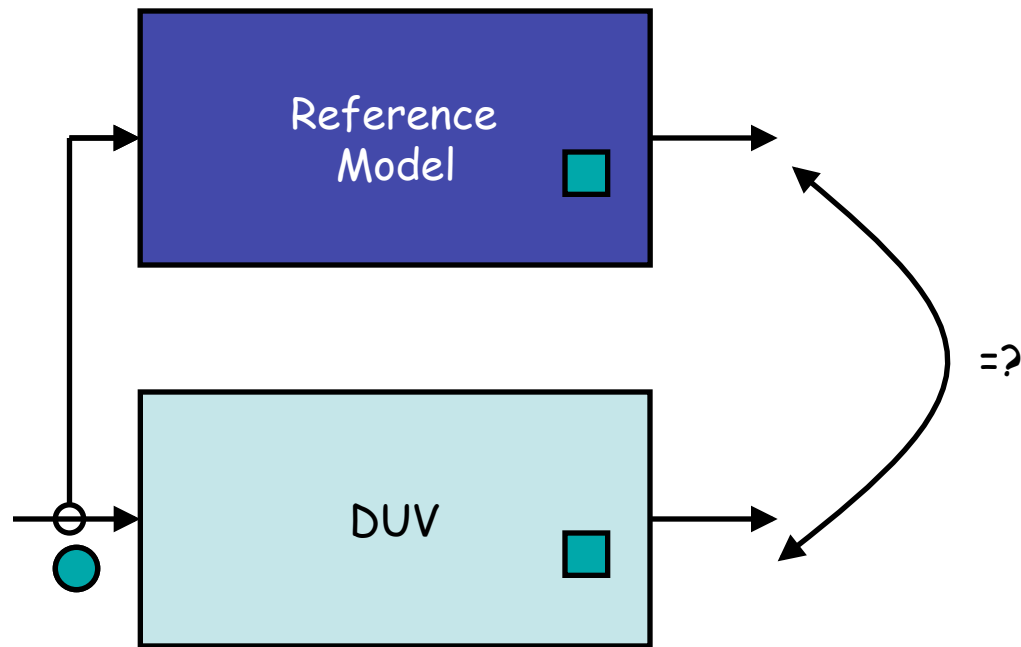
Side Note – Graceful End-of-test

- Checking that nothing is lost is very important
- If an input does not have a matching output, how can we distinguish between these two cases:
 - The input is lost or hopelessly stuck in the DUV
 - The DUV did not have enough time to handle the input
- Possible solution – Start a timer when a new input enters the DUV
 - If the timer expires, that input is lost or stuck
 - But, what if the delay cannot be bound?
- Alternative (or complementary) solution – stop the inputs before the end of the test and let the design clean itself
 - Because there are no new inputs, things that are stuck inside have a chance to get free

Reference Models

- A reference model is an oracle that tells how the DUV should behave
 - Usually in the form of an alternative implementation
- It runs in parallel to the DUV, using the same inputs and provides the checking mechanisms with information about the expected behavior
 - Checking is done by comparing the expected behavior to the actual one
- Pure reference models can run independently of the DUV
 - But not all reference models are pure (example later)

Reference Model Operation



Reference Models

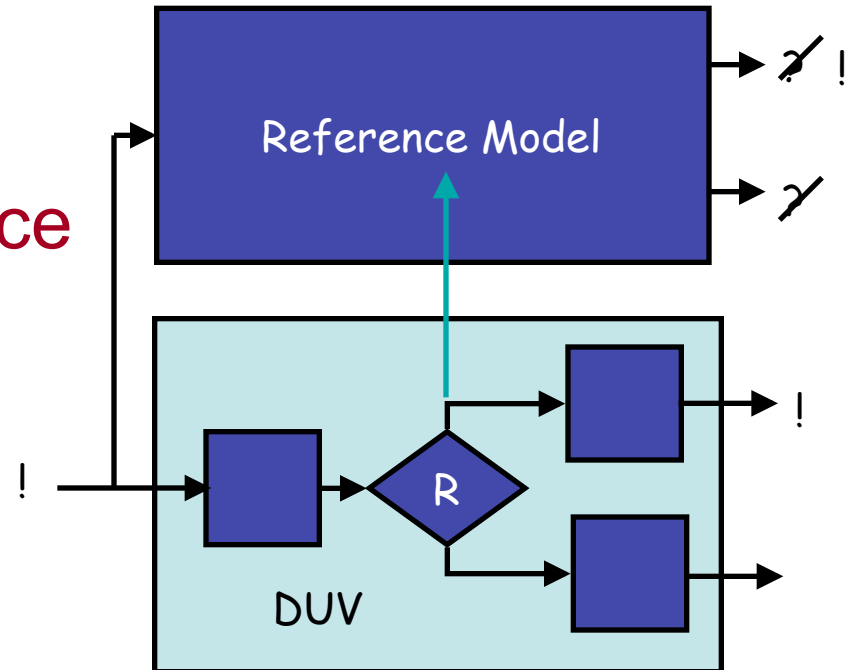
- Reference models have many uses
 - Checking
 - Aids for stimuli generation *(When?)*
 - “Smart” protocol models – imitate the function of the DUV
 - Vehicles for SW development
- What can we check with a reference model
 - In principle, anything
 - In practice it depends on the level of detail and accuracy of the reference model
 - And how much of its behavior we are willing to expose

Levels of Abstraction

- The level of abstraction in a reference model dictates the type of information we can get out of it for checking
 - **Functionally accurate models** can be used only to check correctness of data, usually at the end of the test or at well defined points in time
 - Timing, order, and other checks need other means
 - **Cycle accurate models** can be used for checking all aspects of I/O behavior
 - **Cycle accurate and latch accurate models** can be used also for checking the internal state of the DUV
 - This type of model is sometimes called deep function reference model

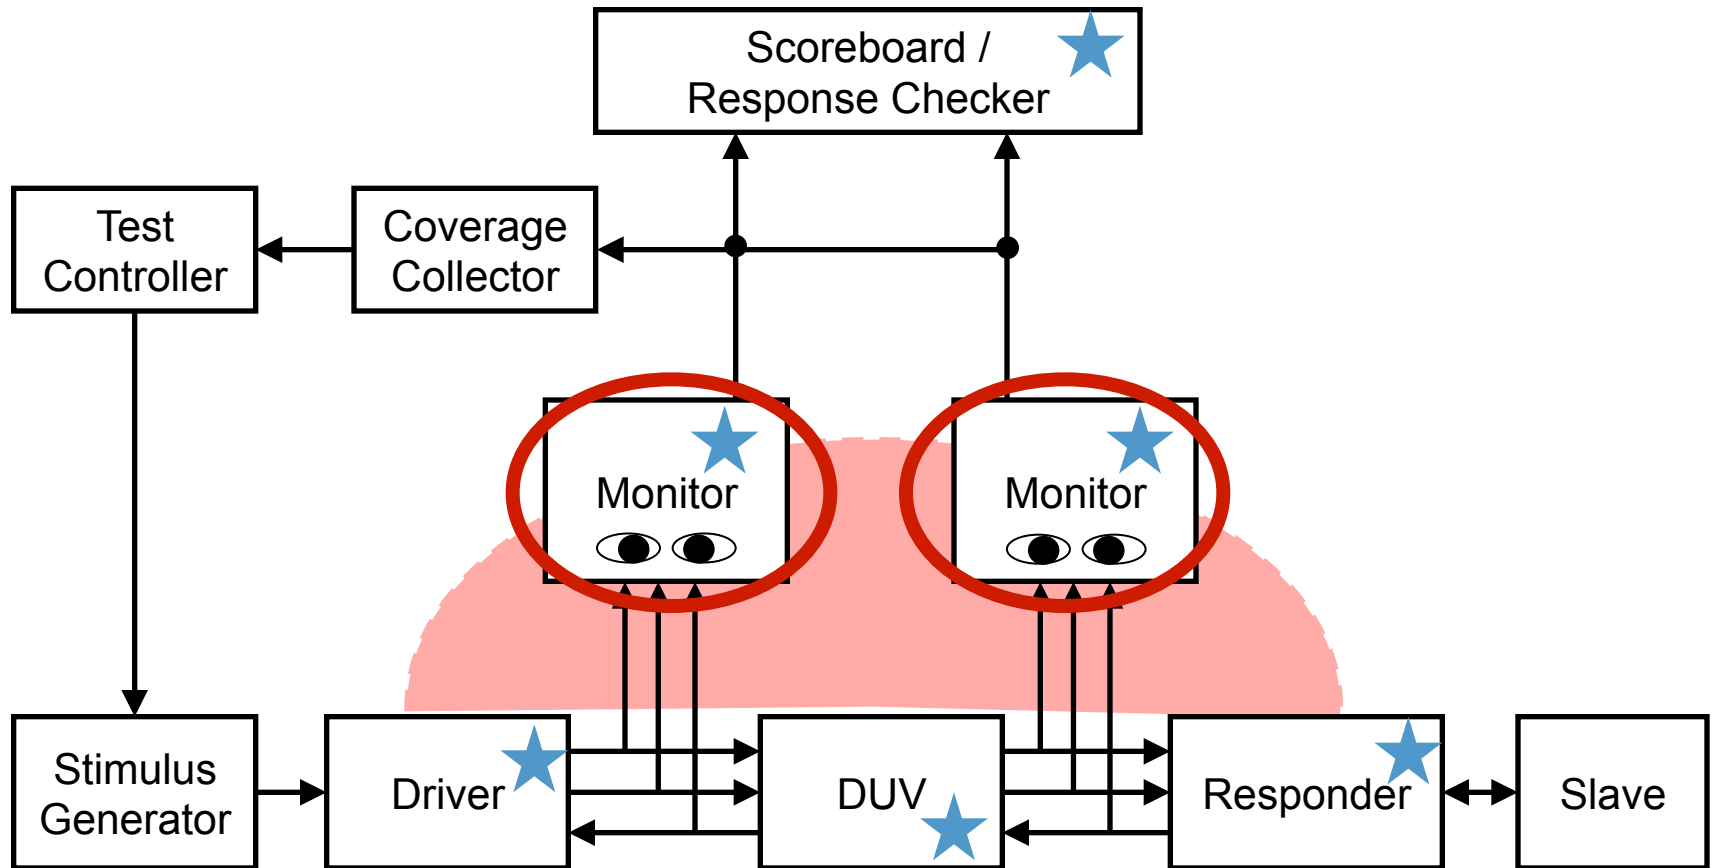
Impure Reference Model

- Sometimes it is impossible (or very hard) for the reference model to duplicate significant decisions made by the DUV
- Possible solution:
Use information from the DUV to assist the reference model!



Contemporary TB Architecture

Contemporary Testbench Architecture



★ Assertions

Monitors

- Monitors are TB components that observe the inputs, outputs, or internals of the DUV.
 - Monitors watch activity of the DUV.
 - Black box: DUV inputs and outputs
 - Grey box: potentially selected internals
 - Monitors can convert low-level signals to transactions.
 - Monitors can flag simple timing and protocol errors.
 - Monitors collect functional coverage.
 - Monitors update the scoreboard.
 - Monitors don't drive DUV pins; they are “passive”.
 - Monitors are self-contained and don't cause “side effects”.
 - Monitors are re-usable at different levels of abstraction.

Types of Monitors

- Input monitors:
 - Collect inputs to the DUV and pass them to scoreboard.
 - Can have checker components.
- Output monitors:
 - Observe the outputs from the DUV and pass them to the scoreboard.
 - Can have checker components.
- Coverage monitors:
 - Collects inputs, outputs and selected internal signals.
 - Permit analysis of stimulus and functionality coverage.

Rule-based Checking

- Checks that a set of rules hold in the DUV
- Essentially, all checking is rule-based, e.g.

if (not “something”) then error

- where the “something” can be
 - Value of a register matches value in reference model
 - Data in a packet at the DUV output matches data in the input as stored in the scoreboard
 - `response_out == 0 → data_out == 0`

Rule-based Checking

- Rules can come from many sources
 - All levels of the design process
 - Spec, high-level design, implementation
 - Behavior of neighboring units
- Rules checking can be implemented in many places
 - External checking tools
 - Various places in the verification environment
 - Interface monitors
 - Scoreboards
 - End-of-test checkers
 - In the DUV itself
- Rule-based checkers embedded in the DUV code are called **assertions**
 - Lecture on **Assertion-Based Verification**

Self-Checking Testbenches

- Knowledge of the DUV functionality can be built into the TB.
 - This *automates* the checking process.
 - Verification engineers encode their knowledge of correct DUV functionality into the checkers, monitors and scoreboard using:
 - Golden Vectors,
 - Reference Models,
 - Protocols or Transactions,
 - Assertions.
- This results in a **“self-checking” TB.**
 - Checkers are “always” active.

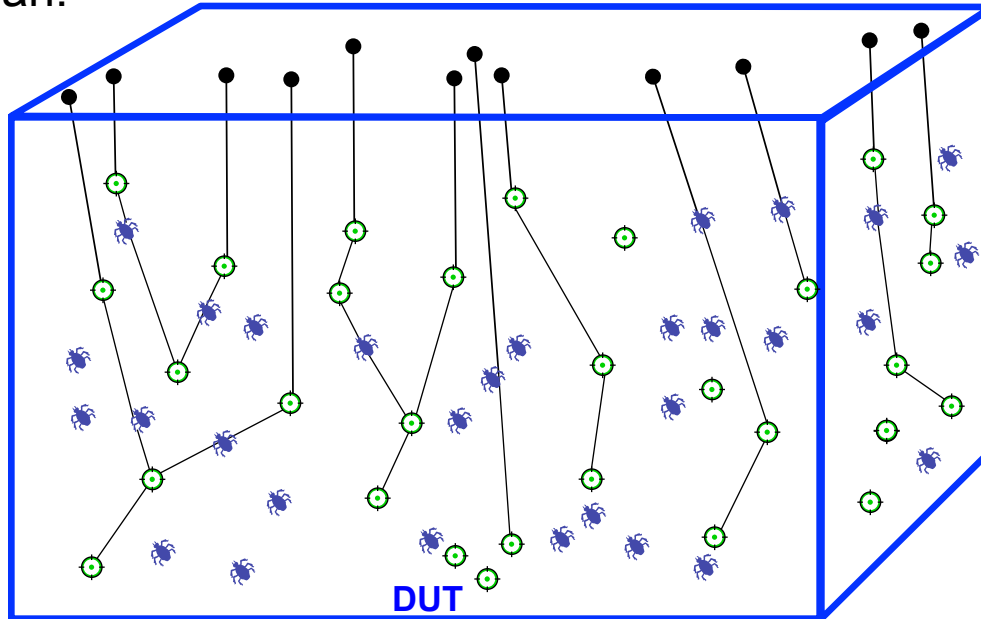
Putting Coverage, Generation and
Checking together:

The Verification Environment

(With many thanks to Cadence for providing the animations in this section.)

Traditional Approach: Directed Testing

Verification engineer sets goals and writes directed test for each item in the Verification Plan:



Redo if design changes

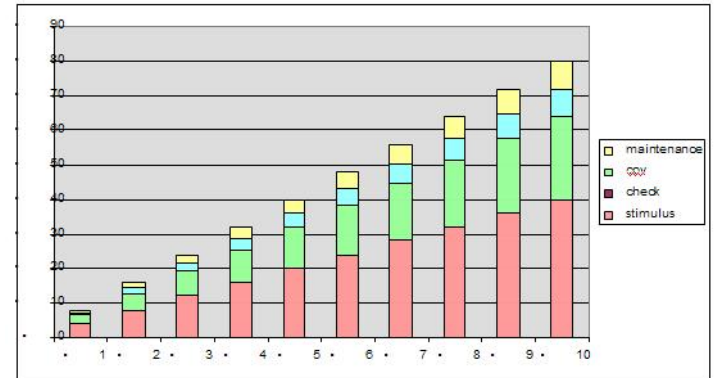
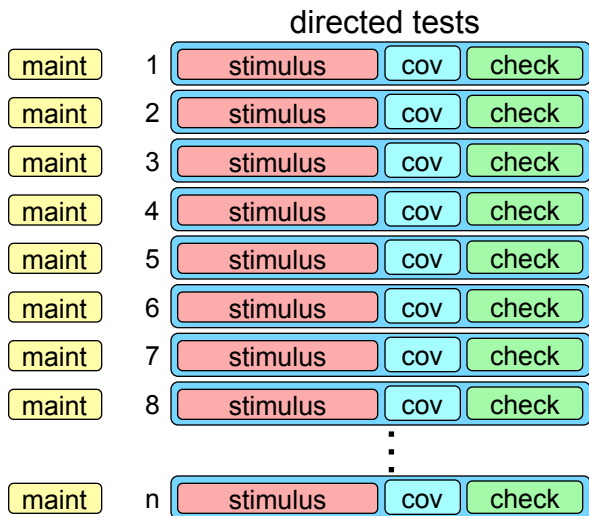
Automation Significant manual effort to write all the tests

Automation Work required to verify each goal was reached

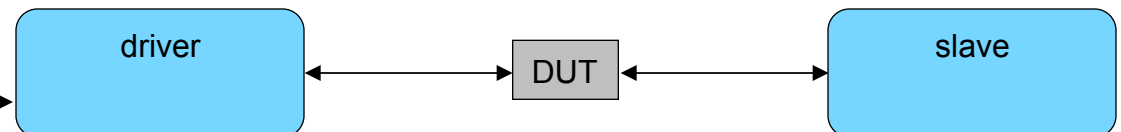
Completeness Poor coverage of non-goal scenarios
... especially the cases that you didn't "think of"

Directed Test Environment

- Composition of a directed test
 - Directed tests contain more than just stimulus.
 - Checks are embedded into the tests to verify correct behavior.
 - The passing of each test is the indicator that a functionality has been exercised.
- Reusability and maintenance
 - Tests can become quite complex and difficult to understand the intent of what functionality is being verified
 - Since the checking is distributed throughout the test suite, it is a lot of maintenance to keep checks updated
 - It is usually difficult or impossible to reuse the tests across projects or from module to system level
- The more tests you have the more effort is required to develop and maintain them

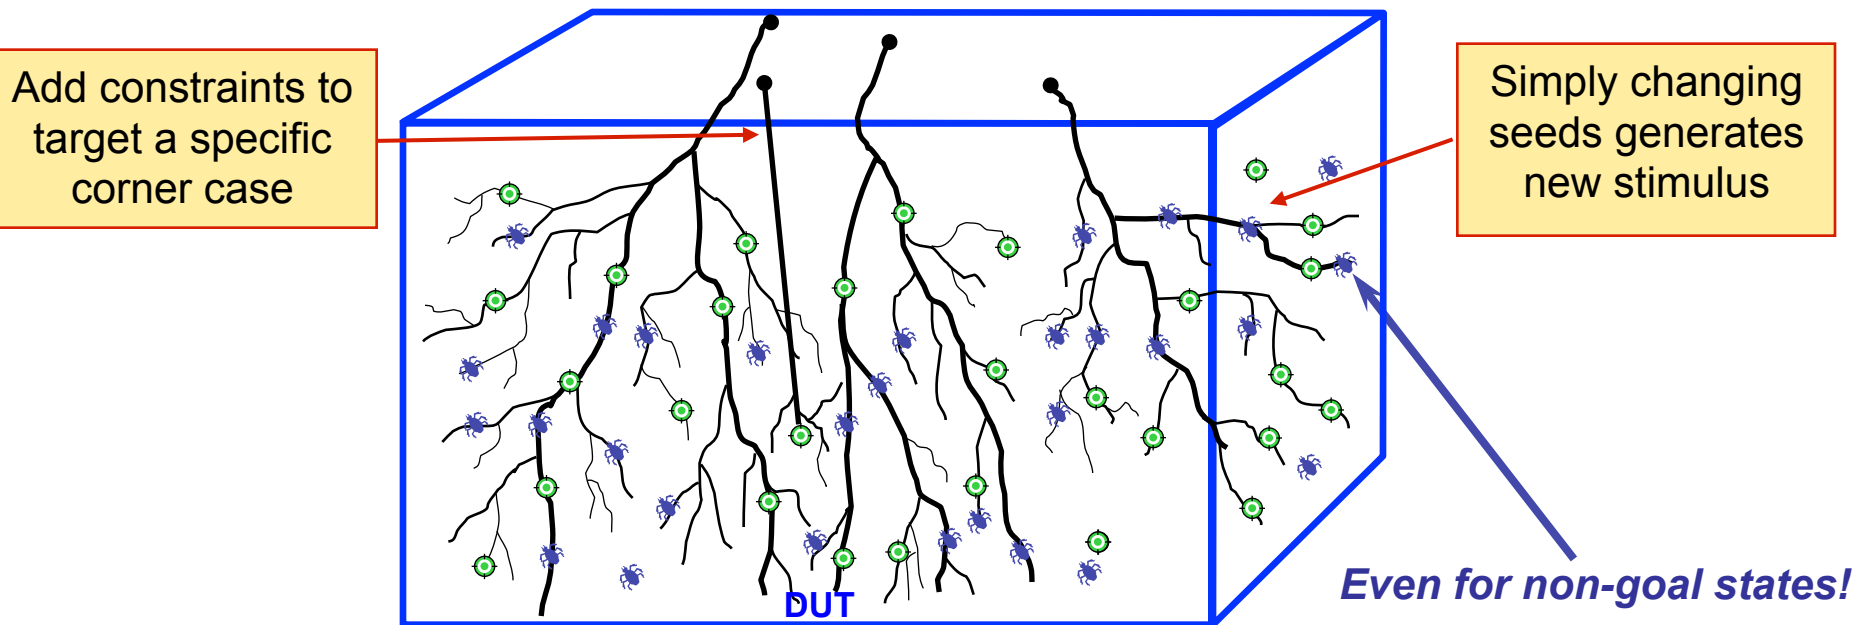


Directed test approach



Coverage Driven Verification Methodology: *Defining Coverage “Goals” Enables Automation*

Focuses on reaching **goal areas** (*versus execution of test lists*):



Constrained-random stimulus **generation** explores goal areas (& beyond).

Coverage shows which **goals** have been exercised and which need attention.

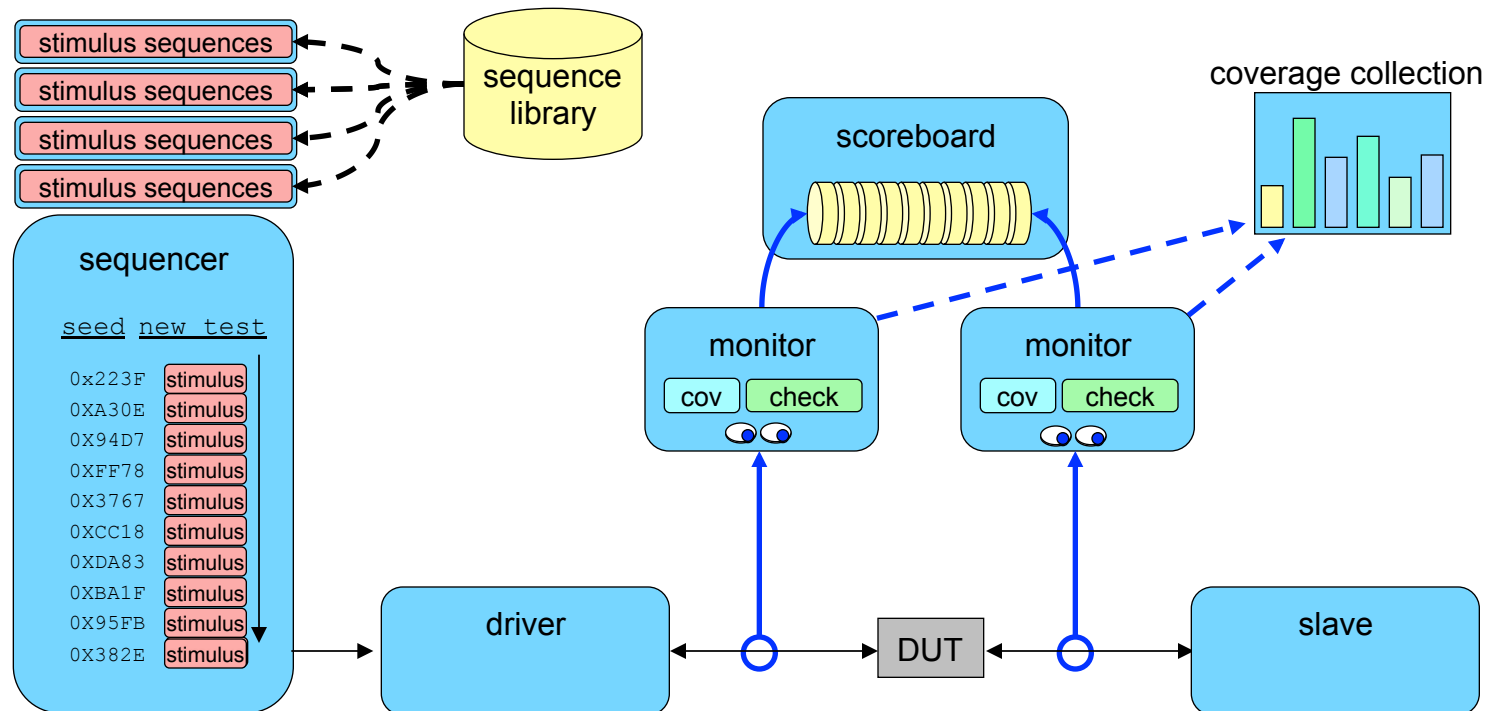
(Checking needs to be in place to assess DUT response.)

Automation – Constrained-random stimulus accelerates hitting coverage goals and exposing bugs. Coverage and checking results indicate effectiveness of each simulation, which enables scaling many parallel runs.

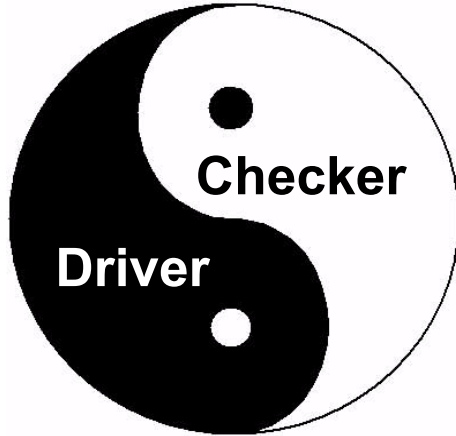
Coverage Driven Environment

■ Composition of a coverage driven environment

- Reusable stimulus sequences developed with “constrained random” generation
- Running unique seeds allows the environment to exercise different functionality
- Monitors independently watch the environment
- Independent checks ensure correct behavior.
- Independent coverage points indicate which functionality has been exercised.



Summary



- Stimuli Generation
- Coverage and
- Checking

Coverage Driven Verification
Methodology

→ Coverage Directed Test Generation